

## **Accelerating CUDA C++ Applications with Multiple GPUs**

Note: The following is not a comprehensive list of topics. You may be asked questions on topics not included below to show your mastery of the subject matter.

In order to become certified to instruct DLI's workshop, you must be prepared to explain the following concepts / terminology:

- JupyterLab
- Basic knowledge of C++ classes and objects
- CUDA thread hierarchy
- CUDA kernels
- Concurrent CUDA streams and their behaviors
- Memory management and data transfers
- Data chunking strategy for multi-stream and multi-GPU
- Copy/compute overlap with multiple GPUs
- CUDA error handling
- NVIDIA®Nsight™ Systems Visual Profiler

The format for demonstrating knowledge in these areas will be a question and answer session with the expectation that you can demonstrate the ability to field student questions and understand the concepts at a deeper level than those attending a workshop.